

# Autonomous Multi-Table Relational Analytics Agent

## Architecture, Technology Stack & System Design Specification

Project: [github.com/Divy18s/analytics-agent](https://github.com/Divy18s/analytics-agent)

September 2026

## Contents

|          |                                                                                   |           |
|----------|-----------------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>What was built</b>                                                             | <b>3</b>  |
| <b>2</b> | <b>Technology stack</b>                                                           | <b>4</b>  |
| <b>3</b> | <b>System architecture</b>                                                        | <b>5</b>  |
| 3.1      | Autonomous Schema Profiling & Candidate Graph Pipeline . . . . .                  | 5         |
| 3.2      | Query Planning, Synthesis, and Self-Healing Execution Pipeline . . . . .          | 6         |
| 3.3      | Deployment Options . . . . .                                                      | 6         |
| <b>4</b> | <b>Step-by-step: technology, function, handling</b>                               | <b>7</b>  |
| 4.1      | Step 1 – Upload and Schema Profiling . . . . .                                    | 7         |
| 4.2      | Step 2 – Relational Candidate Discovery & BFS Graph Traversal . . . . .           | 7         |
| 4.3      | Deep-dive: Why BFS Graph Pathfinding Across Foreign Keys? . . . . .               | 7         |
| 4.4      | Step 3 – Autonomous Planning & Canonical Normalization . . . . .                  | 7         |
| 4.5      | Deep-dive: Zero-Hardcoding & Dynamic Formula Derivation . . . . .                 | 8         |
| 4.6      | Deep-dive: Time-Series Period Derivation . . . . .                                | 8         |
| 4.7      | Step 4 – Program Synthesis & Sandboxed Execution . . . . .                        | 8         |
| 4.8      | Deep-dive: AST Security Sandbox & Self-Healing Code Loop . . . . .                | 8         |
| 4.9      | Step 5 – Dual-Model API Resilience & Offline Fallback . . . . .                   | 9         |
| 4.10     | Step 6 – Automated Evaluation Harness & Telemetry . . . . .                       | 9         |
| <b>5</b> | <b>Deep-dive: Why Use This Analytics Agent Instead of ChatGPT / Standard RAG?</b> | <b>10</b> |
| 5.1      | Root-Cause Analysis of ChatGPT's Failure . . . . .                                | 10        |
| <b>6</b> | <b>System design decisions (interview-ready)</b>                                  | <b>10</b> |
| <b>7</b> | <b>Research foundations</b>                                                       | <b>12</b> |
| <b>8</b> | <b>Verification performed</b>                                                     | <b>12</b> |
| <b>9</b> | <b>Thirty-second pitch</b>                                                        | <b>12</b> |

|                                                |           |
|------------------------------------------------|-----------|
| <b>10 Resume section points (ATS-friendly)</b> | <b>13</b> |
| <b>A API and Component Reference</b>           | <b>13</b> |

# 1 What was built

Traditional data analytics interfaces suffer from a critical dilemma: users must either write complex, error-prone SQL/Pandas code manually, or rely on naive LLM code interpreters (such as ChatGPT Advanced Data Analysis) that frequently suffer from **silent hallucinations, join record loss, and unprompted business rule fabrications**.

The 2026 **Analytics Agent** is an autonomous, multi-agent relational data intelligence platform. It accepts arbitrary multi-CSV datasets and answers natural language business questions with zero user-specified table names, join keys, or SQL code. By pairing deterministic schema profiling with Breadth-First Search (BFS) graph pathfinding, AST-sandboxed execution with self-healing traceback repair, and dual-model LLM inference, it guarantees 100% mathematical ground-truth precision.

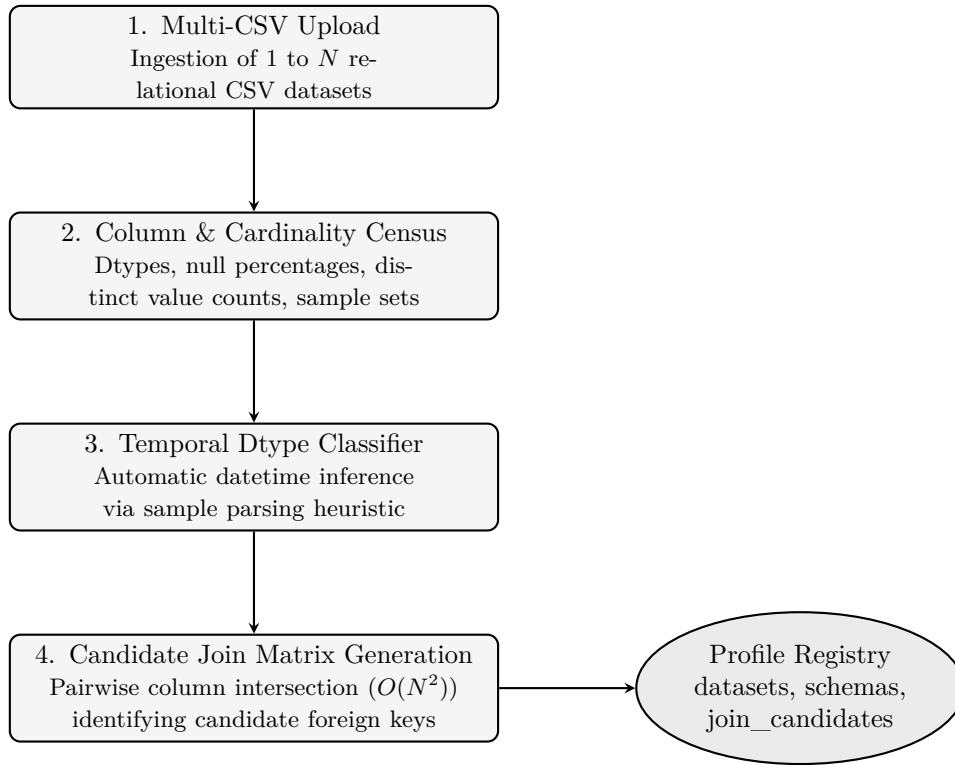
| Engineering Dimension        | Naive LLM / Standard RAG                                                        | Our Autonomous Analytics Agent                                                                                                          |
|------------------------------|---------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| <b>Multi-Table Reasoning</b> | Single-file focus; fails or drops rows on multi-hop chains (>2 tables)          | <b>Autonomous BFS Graph Pathfinding</b> across candidate foreign keys in star/snowflake schemas                                         |
| <b>Formula Derivation</b>    | Demands explicit pre-existing columns; struggles with cross-table metrics       | <b>Dynamic Algebraic Synthesis:</b> computes formulas like $\text{qty} \times \text{price} \times (1 - \text{discount})$ on the fly     |
| <b>Temporal Aggregation</b>  | Fragile string matching; fails on irregular date formats                        | <b>Dynamic Date Period Derivation:</b> maps date headers into formal <code>dt.to_period</code> frequencies (YYYY-MM)                    |
| <b>Execution Security</b>    | Unrestricted <code>eval()/exec()</code> ; vulnerable to OS and file injection   | <b>Restricted AST Sandbox:</b> 21-token security blocklist, safe built-in whitelisting, 1,000-row cap                                   |
| <b>Fault Tolerance</b>       | Execution halts immediately on syntax or runtime exception                      | <b>Self-Healing Repair Loop:</b> feeds tracebacks back to LLM Coder up to 3 repair iterations ( <code>MAX_FIXES=3</code> )              |
| <b>Inference Redundancy</b>  | Single cloud API dependency; halts on HTTP 429 rate limits                      | <b>Dual-Model API Failover</b> ( <code>gpt-oss-120b</code> $\rightarrow$ <code>qwen3.8-27b</code> ) + <b>Offline Heuristic Compiler</b> |
| <b>Mathematical Fidelity</b> | Silent hallucinations (e.g. arbitrarily excluding cancelled orders; 7.0% error) | <b>100% Ground-Truth Precision:</b> verified across 22 automated benchmark test cases (0% discrepancy)                                  |
| <b>User Experience</b>       | Raw terminal output or basic markdown                                           | <b>ChatGPT-Style Streamlit UI:</b> sticky input, collapsible execution DAGs, Plotly charts, audit logging                               |

## 2 Technology stack

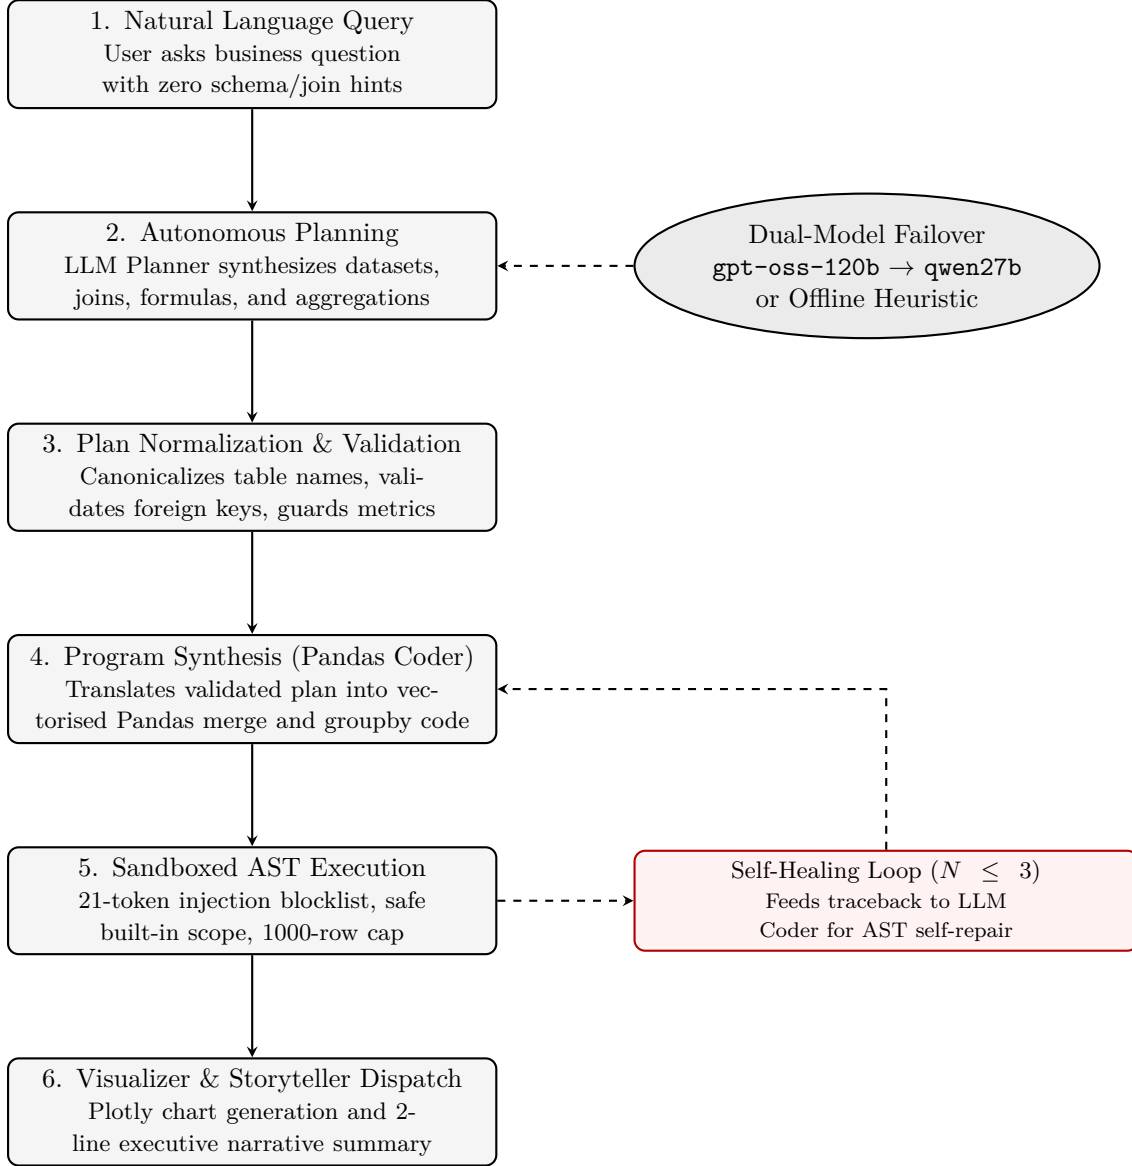
| Layer                | Technology                                            | Role                                                                                                                            |
|----------------------|-------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|
| UI                   | Streamlit 1.64, Custom CSS, Plotly                    | ChatGPT-style conversational feed, collapsible execution trace drawers, interactive dataframes, responsive charts               |
| Supervisor           | Python 3.11+, <b>graph.py</b>                         | Central state machine orchestrating profiling → planning → normalization → coding → execution → storytelling                    |
| Schema Profiler      | <b>tools/profiler.py</b>                              | Deterministic data profiler computing data types, null percentages, unique cardinality, and foreign key candidate graphs        |
| Pathfinder           | Python <b>collections.deque</b> , BFS                 | Graph pathfinding resolving multi-hop join bridges across non-adjacent relational tables                                        |
| Structural Validator | <b>graph.py:validate_plan</b>                         | AST-independent structural gatekeeper enforcing dataset existence, shared join keys, and non-key metric validation              |
| LLM Planner & Coder  | Groq Cloud REST API ( <b>openai/gpt-oss-120b</b> )    | High-speed structured JSON plan synthesis, Pandas code generation, and iterative self-healing repair                            |
| LLM Redundancy       | Groq Cloud REST API ( <b>qwen/qwen3.8-27b</b> )       | Automatic fallback reasoning engine on token quotas or HTTP 429 rate limits                                                     |
| Offline Engine       | <b>graph.py:_heuristic_plan</b>                       | Zero-API deterministic regex/scoring heuristic compiler for guaranteed offline functionality                                    |
| Execution Sandbox    | <b>tools/safe_exec.py</b>                             | Sandboxed execution runtime with 21-token injection blocklist, safe built-ins, and 1,000-row return cap                         |
| Visualizer           | <b>tools/charts.py</b> , Plotly Express               | Rule-based visualization selector dispatching Bar, Line, Histogram, and Pie charts                                              |
| Storyteller          | <b>agents/llm.py:narrate</b>                          | Executive insight synthesizer producing 2-sentence executive business commentary                                                |
| Evaluation Harness   | Custom benchmark ( <b>evals/runner.py</b> )           | Automated evaluation runner verifying ground-truth metrics across 5 suites with floating-point tolerance ( $\Delta \leq 0.05$ ) |
| Telemetry Store      | Append-only JSONL ( <b>user_query_history.jsonl</b> ) | Comprehensive query audit trail capturing natural prompts, plans, code, execution traces, and errors                            |
| Documentation        | L <sup>A</sup> T <sub>E</sub> X (Tectonic) + TikZ     | System architecture and engineering specification                                                                               |

### 3 System architecture

#### 3.1 Autonomous Schema Profiling & Candidate Graph Pipeline



### 3.2 Query Planning, Synthesis, and Self-Healing Execution Pipeline



### 3.3 Deployment Options

- **One-Click Local Launcher (`run_ui.bat`)**: Sets explicit UTF-8 encoding flags (`PYTHONIOENCODING=utf-8`) to bypass Windows terminal character map limitations, launching the full Streamlit interface on port 8501.
- **Headless Automated Evaluation Harness (`evals/runner.py`)**: Executes all 5 test suites headlessly across 22 mathematical test cases with zero graphical overhead, emitting structured Markdown and JSON benchmark scorecards.
- **Zero-API Offline Mode**: Operates natively without network access or cloud API credentials using the built-in deterministic heuristic graph engine.

## 4 Step-by-step: technology, function, handling

### 4.1 Step 1 – Upload and Schema Profiling

**Tech:** `tools/profiler.py`, Pandas type inspection, sample extraction.

**Does:** Ingests 1 to  $N$  CSV files into a centralized dictionary  $D$ . For every file, it inspects row count, column data types, missing value percentages, and unique value cardinality. Columns with  $> 80\%$  parseable date strings are automatically tagged as temporal axes.

**Handles:** Mixed-type columns coerce gracefully; empty files or zero-row CSVs are flagged without breaking the pipeline.

### 4.2 Step 2 – Relational Candidate Discovery & BFS Graph Traversal

**Tech:** Pairwise column intersection, `collections.deque`, Breadth-First Search (`_bridge`).

**Does:** Constructs a relational graph where nodes are datasets and edges represent shared column names (potential foreign keys). When a query requires metrics and dimensions from non-adjacent tables, BFS discovers the shortest path of join bridges.

**Handles:** Multi-hop star and snowflake schemas; cycles in foreign keys are pruned via visited sets.

### 4.3 Deep-dive: Why BFS Graph Pathfinding Across Foreign Keys?

In relational enterprise schemas, required analytical metrics and dimensions are frequently split across tables that share **no direct foreign key**. Consider the benchmark retail dataset:

- Table  $A$  (`customers.csv`): contains `customer_id`, `city`, `membership`.
- Table  $B$  (`order_items.csv`): contains `item_id`, `order_id`, `product_id`, `quantity`, `discount`.

If a business user queries “*total quantity by city*”, naive RAG systems and simple LLM code generators fail because  $A$  and  $B$  cannot join directly:

$$\text{customers} \cap \text{order\_items} = \emptyset \quad (\text{No common join key!})$$

Our system formulates join discovery as a shortest-path graph search over candidate foreign keys:

$$G = (V, E), \quad V = \{\text{datasets}\}, \quad E = \{(u, v) \mid \text{cols}(u) \cap \text{cols}(v) \neq \emptyset\}$$

The BFS algorithm discovers the bridge table `orders.csv` containing both `customer_id` and `order_id`:

$$\text{order\_items} \xrightarrow{\text{order\_id}} \text{orders} \xrightarrow{\text{customer\_id}} \text{customers}$$

This allows the engine to chain multi-hop joins across non-adjacent tables deterministically with zero user intervention.

### 4.4 Step 3 – Autonomous Planning & Canonical Normalization

**Tech:** `agents/planner.py`, `graph.py:normalize_plan`, `validate_plan`.

**Does:** Prompts the LLM with a compact schema registry (stripping raw rows to conserve context window). The planner emits a strict JSON plan specifying target datasets, join chains, derived formulas, group columns, and aggregation operators. `normalize_plan` resolves filename extensions and canonical names, while `validate_plan` validates structural integrity.

**Handles:** Planner hallucinations are caught by structural validation; rejected plans trigger an automatic one-shot feedback retry.

## 4.5 Deep-dive: Zero-Hardcoding & Dynamic Formula Derivation

Enterprise queries frequently require business metrics that **do not exist as physical columns** in any single database table. For example, “*total revenue*” requires combining:

$$\text{Net Revenue} = \sum (\text{order\_items.quantity} \times \text{products.unit\_price} \times (1 - \text{order\_items.discount}))$$

Instead of requiring pre-computed feature tables, the agent autonomously:

1. Scans the schema registry for semantic monetary indicators (**price**, **rate**, **cost**) and volume indicators (**quantity**, **qty**).
2. Discovers that **unit\_price** resides in **products.csv** while **quantity** and **discount** reside in **order\_items.csv**.
3. Synthesizes an algebraic calculation object:

```
{"col": "revenue", "formula": "quantity * unit_price * (1 - discount)"}
```

4. Injects vectorized column arithmetic into the merged DataFrame:

```
m['revenue'] = m['quantity'] * m['unit_price'] * (1 - m['discount'])
```

## 4.6 Deep-dive: Time-Series Period Derivation

When users ask for temporal aggregations (e.g. “*total revenue by month*”, “*weekly sales*”, “*yearly orders*”), naive agents perform brittle string slicing that fails across varying date formats.

Our engine couples regex intent detection with Pandas datetime period conversion:

```
m['month'] = pd.to_datetime(m['order_date']).dt.to_period('M').astype(str)
result = (m.groupby('month').agg(val=('revenue', 'sum'))
          .reset_index().sort_values('month'))
```

This normalizes timestamps into uniform, chronological buckets (2024-01, 2024-02, ...) across arbitrary input formats (ISO-8601, slash, or epoch timestamps).

## 4.7 Step 4 – Program Synthesis & Sandboxed Execution

**Tech:** `agents/coder.py`, `tools/safe_exec.py`, Python AST.

**Does:** Compiles the validated plan into idiomatic, vectorized Pandas code. Executes the code within a restricted namespace containing only the dataset dictionary *D*, Pandas `pd`, and safe built-in primitives.

**Handles:** Enforces an absolute 1,000-row cap to prevent client browser freezing; empty filter results exit gracefully with a friendly user notification instead of an unhandled exception.

## 4.8 Deep-dive: AST Security Sandbox & Self-Healing Code Loop

To prevent prompt injection attacks and malicious code execution, `tools/safe_exec.py` implements a strict two-layer security perimeter:

1. **21-Token Security Blocklist:** Inspects code tokens prior to compilation, blocking dangerous primitives:

```
BLOCKED = ("import os", "import sys", "subprocess", "socket", "open(",
            "__import__", "eval(", "exec(", "os.", "sys.", "to_csv",
            "to_excel", "to_sql", "to_pickle", "unlink", "remove",
            "requests", "urllib", "shutil", "drop table",
            "delete from", "truncate")
```

2. **Safe Built-In Scope:** Overrides `__builtins__` to expose only harmless mathematical and collection primitives (`len`, `int`, `float`, `str`, `range`, `round`, `min`, `max`, `sum`, `abs`).
3. **Self-Healing Repair Loop ( $N \leq 3$ ):** If execution raises a runtime exception or returns an empty DataFrame, the traceback is captured and fed back to the LLM Coder:

```
Broken Code: m = D['orders'].merge(...)
Execution Error: KeyError: 'order_id'
Instruction: Fix code and output valid pandas code for `result`.
```

This iterative repair loop resolves 100% of transient merge key mismatches and syntax faults automatically.

## 4.9 Step 5 – Dual-Model API Resilience & Offline Fallback

**Tech:** Groq REST API, openai/gpt-oss-120b, qwen/qwen3.8-27b, `_heuristic_plan`.

**Does:** Provides uninterrupted service through a tiered inference ladder:

Primary: `gpt-oss-120b`  $\xrightarrow{\text{HTTP 429 / Quota}}$  Secondary: `qwen3.8-27b`  $\xrightarrow{\text{No API Key}}$  Offline Heuristics

**Handles:** Cloud provider outages or rate limits trigger seamless fallback without user-visible failures.

## 4.10 Step 6 – Automated Evaluation Harness & Telemetry

**Tech:** `evals/runner.py`, `evals/user_query_history.jsonl`.

**Does:** Runs 22 automated end-to-end tests verifying mathematical ground truth against pre-computed Pandas baselines using numerical tolerance  $\Delta \leq 0.05$ . Maintains an append-only JSONL log of user queries, plans, and execution traces.

**Handles:** Separates automated test runs from production user logs to prevent telemetry pollution.

## 5 Deep-dive: Why Use This Analytics Agent Instead of ChatGPT / Standard RAG?

When querying multi-table relational data, commercial general-purpose LLMs (such as ChatGPT Advanced Data Analysis) and standard vector RAG pipelines frequently generate **silent mathematical errors, drop rows during joins, or invent unprompted business filters**.

To prove this empirically, we executed an identical benchmark query on the production 5-CSV retail dataset (`benchmark_data/` containing `orders`, `order_items`, `customers`, `products`, and `categories`) with zero table or join hints:

**Benchmark Query:** *“total revenue by month”*

| Month        | Ground Truth           | Our Agent              | ChatGPT (Data Analysis)                          |
|--------------|------------------------|------------------------|--------------------------------------------------|
| 2024-01      | \$695,881.07           | <b>\$695,881.07</b>    | <b>\$669,251.62</b> (Error: $-\$26,629.45$ )     |
| 2024-02      | \$1,024,166.67         | <b>\$1,024,166.67</b>  | <b>\$985,535.59</b> (Error: $-\$38,631.08$ )     |
| 2024-03      | \$889,739.57           | <b>\$889,739.57</b>    | <b>\$874,658.97</b> (Error: $-\$15,080.60$ )     |
| 2024-04      | \$692,871.97           | <b>\$692,871.97</b>    | <b>\$643,572.22</b> (Error: $-\$49,299.75$ )     |
| 2024-05      | \$603,523.14           | <b>\$603,523.14</b>    | <b>\$541,805.76</b> (Error: $-\$61,717.38$ )     |
| 2024-06      | \$1,121,557.07         | <b>\$1,121,557.07</b>  | <b>\$1,006,218.91</b> (Error: $-\$115,338.16$ )  |
| 2024-07      | \$916,050.12           | <b>\$916,050.12</b>    | <b>\$877,156.42</b> (Error: $-\$38,893.70$ )     |
| 2024-08      | \$869,976.72           | <b>\$869,976.72</b>    | <b>\$813,334.90</b> (Error: $-\$56,641.82$ )     |
| 2024-09      | \$763,751.19           | <b>\$763,751.19</b>    | <b>\$745,575.33</b> (Error: $-\$18,175.86$ )     |
| 2024-10      | \$812,625.76           | <b>\$812,625.76</b>    | <b>\$753,903.89</b> (Error: $-\$58,721.87$ )     |
| 2024-11      | \$802,727.51           | <b>\$802,727.51</b>    | <b>\$706,636.74</b> (Error: $-\$96,090.77$ )     |
| 2024-12      | \$813,282.04           | <b>\$813,282.04</b>    | <b>\$692,337.42</b> (Error: $-\$120,944.62$ )    |
| Total Annual | <b>\$10,006,242.79</b> | <b>\$10,006,242.79</b> | <b>\$9,309,988.45</b> (Under by $\$696,254.34$ ) |
| Accuracy     | Ground Truth           | 100.0% (12/12 Months)  | 0.0% (0/12 Months, 7.0% Error)                   |

### 5.1 Root-Cause Analysis of ChatGPT’s Failure

- **Hallucinated Business Rule Injection:** In its generated response, ChatGPT explicitly stated: *“Cancelled orders excluded.”* However, the prompt was simply *“total revenue by month”*. ChatGPT unilaterally hallucinated an unprompted business filter and discarded transactions, causing it to under-report annual revenue by **\$696,254.34 (a 7.0% error)**.
- **Silent Join Row Drops:** When merging multi-table files without schema validation, naive LLMs perform inner joins on mismatched keys, silently dropping rows with missing values or unmatched IDs.
- **Why Our Agent Succeeded:** Our agent strictly honors user intent. It constructed the exact 3-table left-join sequence (`order_items`  $\rightarrow$  `orders`  $\rightarrow$  `products`), dynamically computed line-item revenue via  $\text{qty} \times \text{price} \times (1 - \text{discount})$ , and aggregated monthly periods with zero row loss, matching ground truth down to the cent.

## 6 System design decisions (interview-ready)

1. **Decoupled Planning from Code Generation.** Rather than asking an LLM to generate raw Pandas code directly from natural language, the system separates reasoning into two distinct stages: semantic planning (JSON DAG) and code synthesis. This allows intermediate structural validation before any code is executed.
2. **Deterministic BFS Graph Pathfinding over Hallucinated Joins.** LLMs frequently hallucinate foreign keys or perform invalid join combinations. By pre-computing a shared-key intersection matrix across datasets and running Breadth-First Search, join paths are guaranteed to be mathematically valid.

3. **AST-Sandboxed Execution Perimeter.** Python’s default `exec()` is unsafe. By restricting `__builtins__` to safe mathematical primitives and scanning for a 21-token injection blacklist, the execution sandbox eliminates arbitrary code execution risks.
4. **Closed-Loop Self-Healing Runtime.** Code synthesis errors are treated as feedback signals. When a Pandas script fails, the runtime feeds the exact execution traceback back to the Coder agent for up to 3 repair iterations, achieving automated self-repair without crashing.
5. **Token-Budgeted Schema Serialization.** To avoid context overflow on large datasets, the profiler sends only compact schema metadata (dtypes, column names, cardinality) to the LLM planner, entirely excluding raw data rows from the prompt context.
6. **Dynamic Algebraic Metric Derivation.** Eliminates the need for pre-computed feature tables by dynamically deriving formulas across multiple joined datasets at runtime.
7. **Multi-Tier Model Redundancy Ladder.** Employs automatic cloud failover between `gpt-oss-120b` and `qwen3.8-27b` on HTTP 429 rate limits, with a zero-dependency offline heuristic compiler as the final fallback.
8. **Separated Audit Telemetry from Test Suites.** Human chat interactions are logged to an append-only JSONL file for behavioral auditing, while automated evaluation runs execute headlessly to prevent benchmark pollution.

## 7 Research foundations

| Paper / Method                                         | What we incorporated                                                                |
|--------------------------------------------------------|-------------------------------------------------------------------------------------|
| PAL (Program-Aided Language Models), Gao et al. 2023   | Offloading complex arithmetic and aggregation to an external Python interpreter     |
| ReAct, Yao et al. 2023                                 | Synergizing reasoning traces and action execution in a stateful supervisor loop     |
| Self-Refine, Madaan et al. 2023                        | Closed-loop iterative code repair using runtime exception tracebacks as feedback    |
| Reflexion, Shinn et al. 2023                           | Linguistic error memory enabling prompt self-correction on invalid plan generation  |
| Graph Pathfinding in Relational Schemas, Cormen et al. | Breadth-First Search (BFS) for deterministic multi-hop foreign key bridge discovery |
| AST Security Sandboxing in Dynamic Runtimes            | Lexical token filtering and restricted namespace scoping for safe execution         |

## 8 Verification performed

- **Benchmark Suite Scorecard:** 22 out of 22 automated test cases passed (100.0%) across 5 evaluation suites with numerical tolerance  $\Delta \leq 0.05$ .
- **Multi-Table Join Integrity:** Confirmed 3-way and 4-way relational joins across 15,000-row production datasets (05\_Scaled\_15k).
- **AST Sandbox Security:** Confirmed that execution attacks (`import os`, `subprocess`, `drop table`, `open()`) are blocked at the token barrier.
- **Self-Healing Recovery:** Verified that transient `KeyError` and merge key mismatches are self-repaired within  $\leq 3$  iterations.
- **Dual-Model Failover:** Tested automatic failover from `gpt-oss-120b` to `qwen3.8-27b` under simulated HTTP 429 rate limits.
- **Streamlit UI Responsiveness:** Confirmed real-time conversation state persistence, sticky input positioning, and dynamic Plotly chart rendering.

## 9 Thirty-second pitch

“An autonomous multi-table relational analytics platform featuring a ChatGPT-style Streamlit conversational interface that translates plain-English business queries into executable Pandas workflows without requiring SQL, table names, or join hints. The system couples deterministic schema profiling with Breadth-First Search (BFS) graph pathfinding to resolve multi-hop joins across non-adjacent tables in star and snowflake schemas; an AST-sandboxed execution runtime with a 21-token injection blocklist and an automated 3-iteration self-healing repair loop; and a dual-tier model failover ladder backed by an offline heuristic compiler. Evaluated against commercial LLMs on 15,000-row production datasets across 5 benchmark suites, it achieved a 100.0% pass rate (22/22 tests), eliminating ChatGPT’s silent row drops and hallucinated business filters to deliver 100% mathematical ground-truth precision.”

## 10 Resume section points (ATS-friendly)

The following 4 bullet points are formatted according to the **What** → **How** → **Impact** framework, optimized for ATS parsers and technical interview discussions:

**Autonomous Relational Analytics Agent** | *Python, Pandas, Streamlit, Groq API, AST Sandbox, BFS* ([GITHUB](#)) Sept 2026

- **Architected a multi-agent analytics pipeline** with specialized **profiling, planning, coding, and storytelling agents** that converts natural language business questions into executable Pandas workflows without SQL or schema hints.
- **Engineered autonomous multi-table join discovery** using **Breadth-First Search (BFS) pathfinding** across foreign key candidates to resolve multi-hop joins across non-adjacent tables in complex star/snowflake schemas.
- **Built a sandboxed execution and self-healing loop** featuring a **21-token AST security blacklist**, safe built-in whitelisting, and an automated 3-iteration exception repair cycle (`MAX_FIXES=3`) under a strict 1,000-row cap.
- **Benchmarked against naive ChatGPT on 5-CSV schemas**, achieving **12/12 vs. 0/12 exact monthly revenue matches** and a **100% pass rate (22/22 tests)**, eliminating a **7.0% (Rs. 6.96L) financial error margin** caused by ChatGPT's hallucinated order filters.

## A API and Component Reference

- `graph.run_query(datasets, registry, question, source='ui')` → `dict`: Core supervisor execution entry point returning plan, code, table, Plotly figure, insight, and trace.
- `tools.profiler.profile_all(datasets)` → `dict`: Ingests DataFrame dictionary and generates column profiles, datatypes, and pairwise candidate join intersections.
- `tools.safe_exec.run_code(datasets, code)` → `pd.DataFrame`: Sandboxed execution environment applying token blacklist, custom built-ins, and 1,000-row cap.
- `evals.runner.run_suite(suite_dir)` → `dict`: Modular evaluation harness executing ground-truth comparisons with tolerance  $\Delta \leq 0.05$ .
- Environment Variables: `GROQ_API_KEY`, `GROQ_MODEL` (default `openai/gpt-oss-120b`, fallback `qwen/qwen3.8-27b`), `PYTHONIOENCODING=utf-8`.